

What is NoWDB?

Tobias Schoofs

April 20, 2023

Overview

- What?
- Why?
- How?
- Vision!
- Market!
- Where are we now?
- What means the name?

What?

$$\text{NoWDB} = (\textit{graph} + \textit{timeseries} + \text{SQL}) \times (\textit{Python} + \textit{Lua})$$

Why Timeseries? (1/2)

- Growing amount of data are timeseries *to some extent*
- Important driver: IoT
- But also: “traditional” industries:
 - IT infrastructure providers,
 - Energy,
 - Manufacturing (“Industry 4.0”),
 - Retail,
 - Fintech,
 - ...
- Timeseries data have special requirements:
 - Fast ingestion ($> 1M/s$)
 - Scalability of queries ($> 1Bmetrics$)
 - Special handling of time dimension
 - Timeseries are less rigid compared to relational data (time points may be lost or duplicated)

Why Timeseries? (2/2)

- Available timeseries databases are too narrow, *i.e. pure timeseries*.
- Examples:
 - InfluxDB
 - OpenTSDB
- There are projects trying to solve this issue, e.g.:

timescale = postgres + timeseries



... but it's still Postgres: somewhat slow

Why Graph?

- Make timeseries more widely applicable!
- Natural candidates: Relational and Graph
- Graph is more flexible:
we can add new edges
without changing the structure of entities
- Graph is less rigid:
edges pointing to nowhere or
pointing to duplicated data are no issue
- Graph may be easier to grasp:
and this way reduce common errors in data
modelling.
(This, however, is pure speculation on my side ...)



The Founder of Graph Theory:
Leonhard Euler

Why SQL?

- ~~Make SQL great again!~~
- SQL is standard.
- SQL is actively developed with new ideas flowing in.
- SQL is known by virtually everybody in the industry.
- SQL is easy to learn (even though tricky to master . . .)
- SQL is designed for integration with other languages and tools.
- SQL is supported by thousands of tools.
- SQL is not bad anyway.
- Well . . .
- Next slide!

Why Python?



- An important application domain is data science.
- There will be demand for tools, libraries, languages supporting
 - Statistics
 - Calculus
 - Numerical analysis
 - Function analysis and function plotting
 - Linear algebra (vectors, matrices, *etc.*)
 - Graph(!)-oriented data structures
 - ...
- Millions of lines of code available in (or through) Python
- Python is **the** data science language.
(There also is R , but R is very limited as general purpose language)
- It is in particular interesting to have Python available on server side.
This will reduce client/server round trips and network traffic.


Why Lua?

- Python has an issue with its interpreter:
 - difficult to run many interpreters (= sessions)
 - in parallel in one process (= NOWDB server).
- In consequence: Python is not good for tasks that frequently run in parallel.
- We therefore need a second language for more down-to-earth, day-to-day DBA jobs.
- Lua appears to be a good candidate because:
 - Lua is well known in the data science community
 - Lua already had its appearance on the database stage (namely in the redis *in-memory* database)
 - Lua is designed as configuration language (and DBA tasks are often configuration tasks)
 - Lua is designed for integration with the C language.




How? Data Modelling


The Logical Fact



The Relational Way



The Graph Way



How? Graph

Create a vertex type

```
create type supplier (  
    id uint primary key,  
    name text ,  
    ...  
);
```

Create another vertex type

```
create type product (  
    id uint primary key,  
    description text ,  
    ...  
);
```

Add an edge
without changing the types

```
create edge supplies (  
    origin supplier ,  
    destin product  
);
```

How? Timeseries

Product type

```
create type product (  
  id uint primary key,  
  description text,  
  ...  
);
```

Client type

```
create type client (  
  id uint primary key,  
  name text,  
  ...  
);
```

Timeseries edge

```
create edge buys (  
  origin client,  
  destin product,  
  timestamp time,  
  quantity float,  
  price float  
);
```

How? Queries (1/2)

Who bought
product 12345?

(This query runs in ms!)

```
select origin
from buys
where destin = 12345
and year(stamp) = 2018
and month(stamp) >= 10;
```

Traditional Join

```
select p.description
from buys (join product as p on destin)
        (join client as c on origin)
where c.name = 'Smith'
and year(stamp) = 2018;
```

Implicit join

```
select destin.description
from buys
where origin.name = 'Smith'
and year(stamp) = 2018;
```

How? Queries (2/2)

Who are the suppliers of products bought by client 'Smith'?

```
select destin.supplies.origin.name  
from buys  
where origin.name = 'Smith'  
and year(stamp) = 2018;
```

Can we use CTE for recursive search on paths?

Well ...

Next Slide!

```
with path (n,e) as (  
  select 1, origin.friend  
  from friend  
  where origin.name = 'Smith'  
  and year(stamp) = 2018  
union  
  select n+1, destin.friend  
  from friend  
  where year(stamp) = 2018  
  and n < 10)  
select count(*)  
  from path  
where origin.name = 'Muller';
```

How? Storage Engine (1/2)

Two storage types:

① Edge Store:

- optimised for fast ingestion
- data packed to favour typical queries
- special handling of timestamps
- small memory footprint
- optional compression (recommended for timeseries edges)

② Vertex Store:

- consistent data (no duplicates)
- column-oriented (*i.e.* optimised for large data types)
- enforce primary key
- in-memory caches to speed up inserts and queries
- optional compression (recommended only for large tables)

How? Storage Engine (2/2)

Four index types:

① Primary Keys (Vertex)

- B⁺Tree
- storing set of *pageids*

② Secondary Vertex Index

- B⁺Tree
- storing set of primary keys

③ Page Index (Edge):

- B⁺Tree
- storing “embedded” B⁺Trees storing pageids
- used for queries and joins on edges

④ Label Index (Edge):

- B⁺Tree
- storing fixed number of pageids
- used to find related edges

Vision!

- **Scalability!**

Reduce hardware and software license cost by allowing more data to be stored on one node, *i.e.* billions of metrics per server.

- **Data Analysis**

Provide server- and client-side Python modules to ease data science.

- **Integration with big-data tools**

Ease big-data by integration with existing tools *e.g.* Kafka for reliable ingestion of data.

- **Integration with distribution engines**

Ease large-scale computation by integration with existing tools *e.g.* SPARK for distributed computing.

- **Integration with visualisation frameworks**

Ease application development by integration with existing GUI frameworks, *e.g.* Grafana.

- **Integration with traditional tools**

Use ODBC/JDBC to make the platform available to traditional tools.

- **Publish/Subscribe Capability**

Provide pub/sub capability to ease the development of monitoring and real-time data analysis.

Market!

- **Complement** existing OLTP infrastructures
- **Compete** with complex big-data and data science platforms providing a low-cost alternative to frameworks with high cost of ownership (e.g.: Hadoop, large OLAP systems, etc.)
- Target industries with “timeserish” data and need of data science (e.g. IoT, retail, fintech, “Industry 4.0”)
- Provide pre-packaged out-of-the-box solutions (e.g. integration with Kafka, Zookeeper, SPARK & Grafana)
- Operate with a mix of open-source and proprietary licenses (e.g. generic modules for free, specific modules for pay)
- Foster a community among customers to grow repository of available solutions and modules
- Foster community of *certified* service providers

Where are we now? (1/2)

- Prototype is available:
 - Client/Server (with clients for Python, C and Go)
 - SQL with many operators and functions:
 - arithmetic and boolean operators and functions,
 - CASE and other conditionals,
 - time and geospatial functions,
 - trigonometry,
 - string operators and functions,
 - ...
 - Python client- and server-side
 - Simple command line client
 - Foundations for grouping, ordering and joins
 - Storage engine needs revision
 - User manual on the way
 - Query planner still rudimentary
 - Todo: security and user management

Where are we now? (2/2)

- We are currently working with the IoT startup **Loka Systems** (<https://loka.systems>) to identify strengths and weaknesses.
 - The Loka database has about 30M metrics
 - It was imported into nowDB in $< 30s$
 - Queries were implemented by means of server-side Python modules
 - The feedback of the Loka team already led to conceptional and technical improvements
- We created a database with data by the World Meteorological Organisation (WMO)
 - with weather data covering the years 1970 – 2018
 - with more than 1B metrics
 - Import takes about 5 minutes
 - Index-based queries run in milliseconds
 - Benchmarks reveal some weaknesses (e.g. grouping)



www.wmo.int

The funny Name

- a reference to time
- a pun on *noSQL* and *newSQL*
- ...and a homage to Norbert Wiener, the father of timeseries analysis for Control & Communication

