

Product Vision

Tobias Schoofs

April 20, 2023

NoWDB

$$\text{NoWDB} = (\textit{graph} + \textit{timeseries} + \text{SQL} + \textit{Lua}) \times \textit{Python}$$

- A data processing platform
- For large amounts of *timeserish* data
- That can be used with standard interfaces (SQL)
- Emphasising data science on client side
- Emphasising big-data on server side

Timeseries (1/2)

- Growing amount of data are timeseries *to some extent*
- Important driver: IoT
- But also: “traditional” industries:
 - IT infrastructure providers,
 - Energy,
 - Manufacturing (“Industry 4.0”),
 - Retail,
 - Fintech,
 - ...
- Timeseries data have special requirements:
 - Fast ingestion ($> 1M/s$)
 - Scalability of queries ($> 1Bmetrics$)
 - Special handling of time dimension
 - Timeseries are less rigid compared to relational data (time points may be lost or duplicated)

Timeseries (2/2)

- Available timeseries databases are too narrow, *i.e. pure timeseries.*
- Examples:
 - InfluxDB
 - OpenTSDB
- There are projects trying to solve this issue, e.g.:

timescale = relational + timeseries




... but it's still Postgres: somewhat slow

Graph

- Make timeseries more widely applicable!
- Natural candidates: Relational and Graph
- Graph is more flexible:
we can add new edges
without changing the structure of entities
- Graph is less rigid:
duplicated edges are no issue



The Founder of Graph Theory:
Leonhard Euler



Provide full power of data-science tools on client side:

- NumPy
- SciPy
- Pandas
- Matplotlib
- ...

Provide ready-made domain-specific packages:

- Statistics
- Timeseries
- Bayes, Support Vector Machines, Clustering, *etc.*
- Geodata
- Retail (e.g. recommendation engine)



High level of integration client & server:

- Server-side support for domain-specific packages,
i.e. support in SQL and server-side scripting (Lua & Python)
- Stream Processing
- Monitoring
- Integration with messaging tools (e.g. Kafka)
- External Interfaces:
 - native (Python, C, Go, *etc.*)
 - REST (for, e.g., Grafana)
 - HTML (through, e.g., Go)
 - ODBC/JDBC

Strategy

- **Complement** existing OLTP infrastructures
- **Compete** with complex big-data and data science platforms
- Provide a **low-cost** alternative to frameworks with high cost of ownership (e.g.: Hadoop, large OLAP systems, etc.)
- Target industries with **timeserish** data and need of data science (e.g. IoT, retail, fintech, “Industry 4.0”)
- Provide pre-packaged out-of-the-box solutions (e.g. integration with Kafka, Zookeeper, SPARK & Grafana)
- Operate with a mix of open-source and proprietary licenses (e.g. generic modules for free, specific modules for pay)